

Lecture

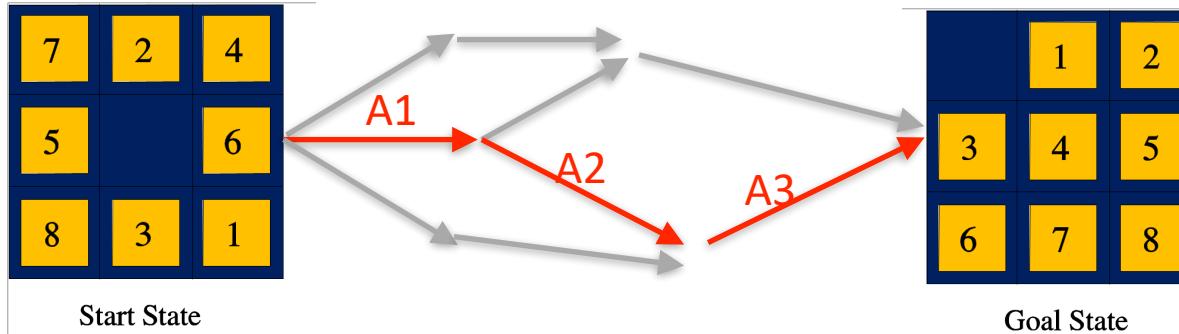
Foundations of Artificial Intelligence

Part 6 – Local Search

Dr. Mohsen Mesgar

Universität Duisburg-Essen

Recall ...



- **Uninformed search**

- BFS: Breadth-first search
- DFS: Depth-first search
- ID: Iterative deepening

- ◊ **Informed search using heuristics**

- ◊ Greedy Best-first search
- ◊ A* search

Any other open questions?



Where are we?

★ Intelligence & Agents

★ Search

 ★ Uninformed

 ★ Informed

 ★ Local

★ Uncertainty & Learning

★ Learning to Represent

Local Search

Optimization Problems

- AI is about setting a goal and optimize the actions to achieve the goal

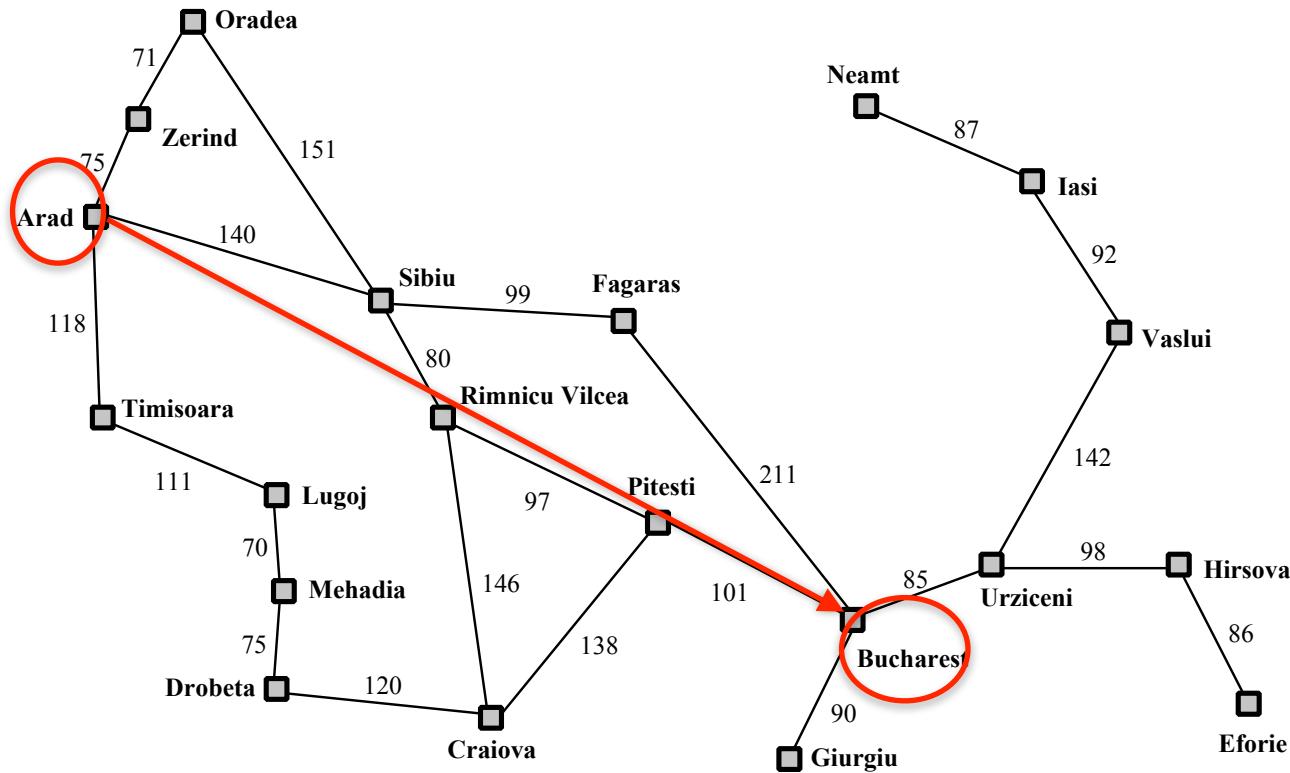
Optimization Problems

- AI is about setting a goal and optimize the actions to achieve the goal
- There are two major optimization problems
 - **Optimal path finding:** start and goal states are clearly defined

Optimization Problems

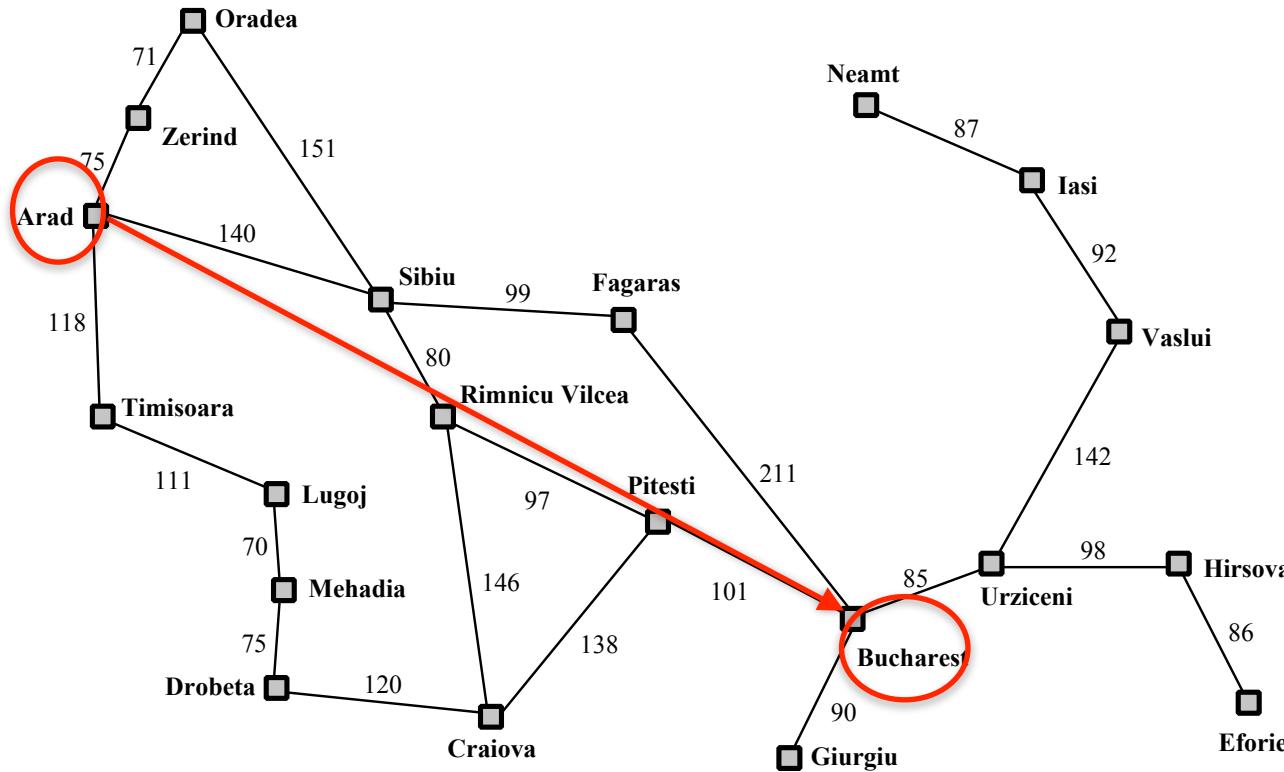
- AI is about setting a goal and optimize the actions to achieve the goal
- There are two major optimization problems
 - **Optimal path finding:** start and goal states are clearly defined
 - **Optimal goal state finding:** start state is given but the goal state is unknown. Instead we **have a set of constraints that should be satisfied**

Path Finding

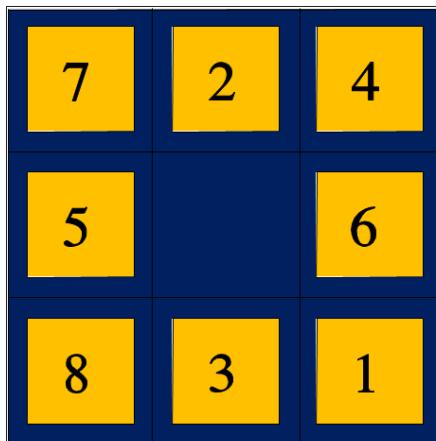


Path Finding

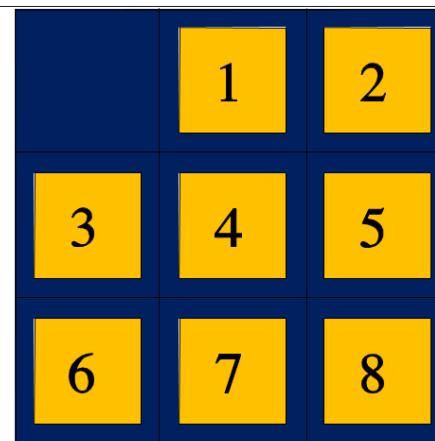
Optimal path finding



8-Puzzle

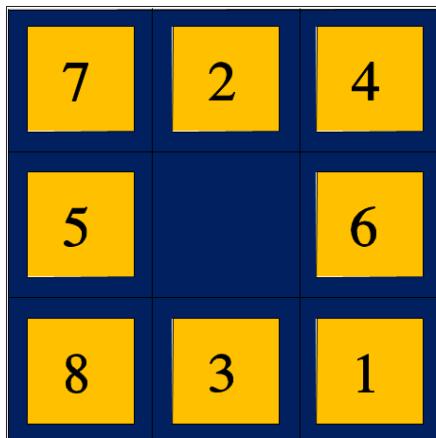


Start State

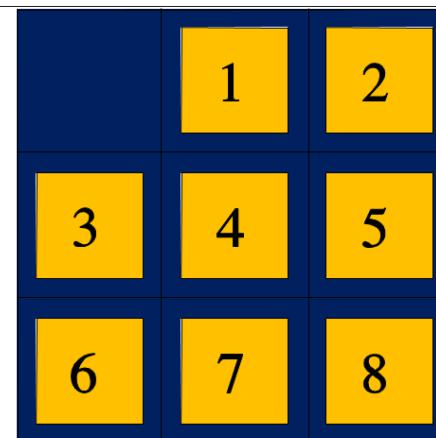


Goal State

Optimal path finding



Start State



Goal State

Sudoku

- **Goal:** To fill a 9×9 grid with digits so that each column, each row, and each of the nine 3×3 subgrids that compose the grid contain all of the digits from 1 to 9

5	3			7				
6			1	9	5			
	9	8				6		
8				6				3
4			8	3			1	
7				2			6	
	6				2	8		
			4	1	9		5	
			8			7	9	

Sudoku

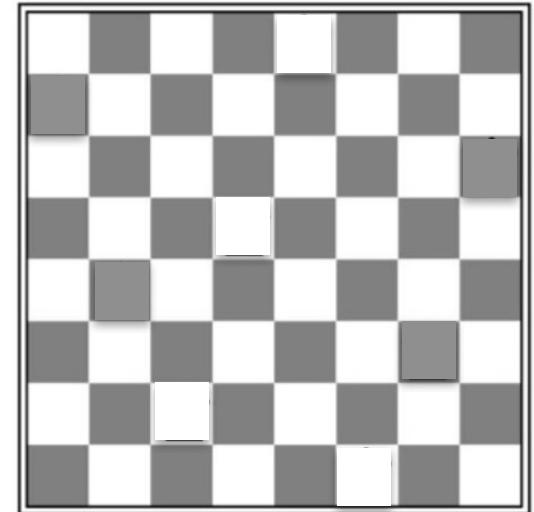
- **Goal:** To fill a 9×9 grid with digits so that each column, each row, and each of the nine 3×3 subgrids that compose the grid contain all of the digits from 1 to 9

Optimal goal state finding

5	3		7					
6			1	9	5			
	9	8				6		
8				6			3	
4			8	3			1	
7				2			6	
	6				2	8		
			4	1	9		5	
				8		7	9	

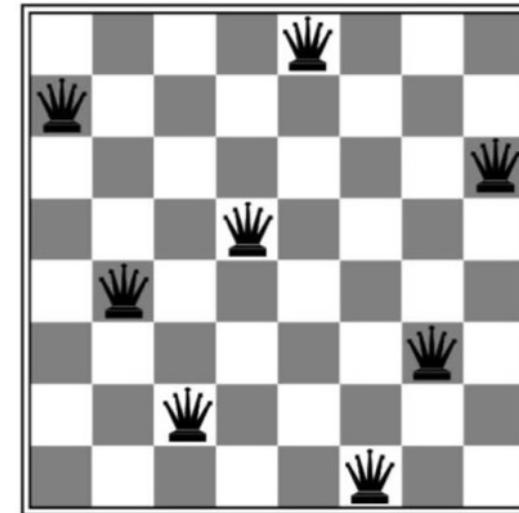
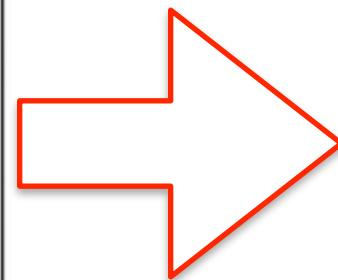
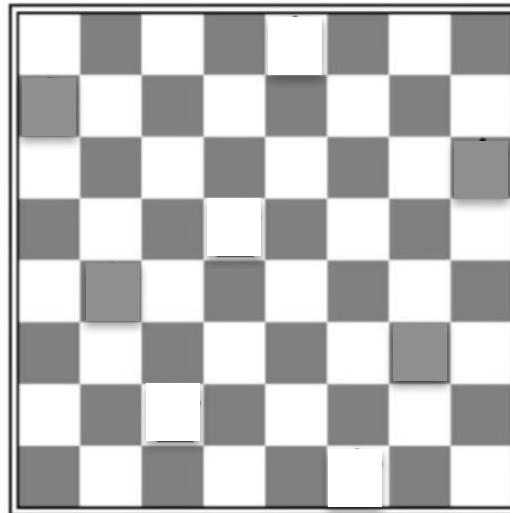
The N-Queens Problem

- **Goal:** To put eight queens on a $N \times N$ chess-board such that none of them threatens any of others.
 - No two queens share the same row,
 - No two queens share the same column
 - No two queens share the same diagonal.



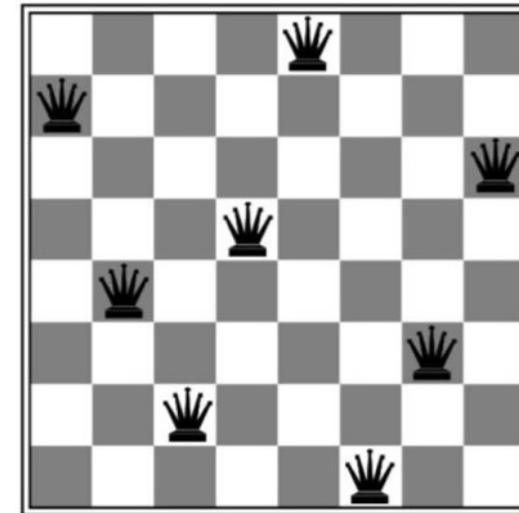
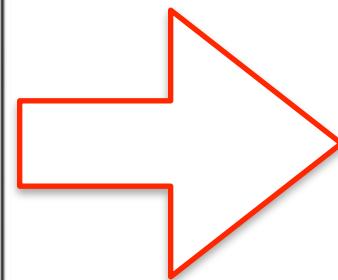
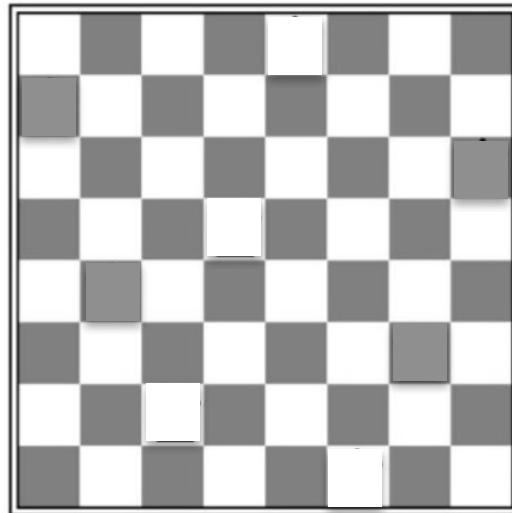
The 8-Queens Problem

- There are 92 goal states. We should find one of them.



Example: The 8-Queens Problem

Optimal goal state finding



Optimization Problems

- There are two major optimization problems
 - **Optimal path finding:** start and goal states are clearly defined
 - Agent uses uninformed or informed search algorithms
 - **Optimal goal state finding:** start state is given but the goal state is unknown. Instead we **have a set of constraints that should be satisfied**
 - Agent uses local search algorithms

State space:

- A set of "complete" configurations

Goal :

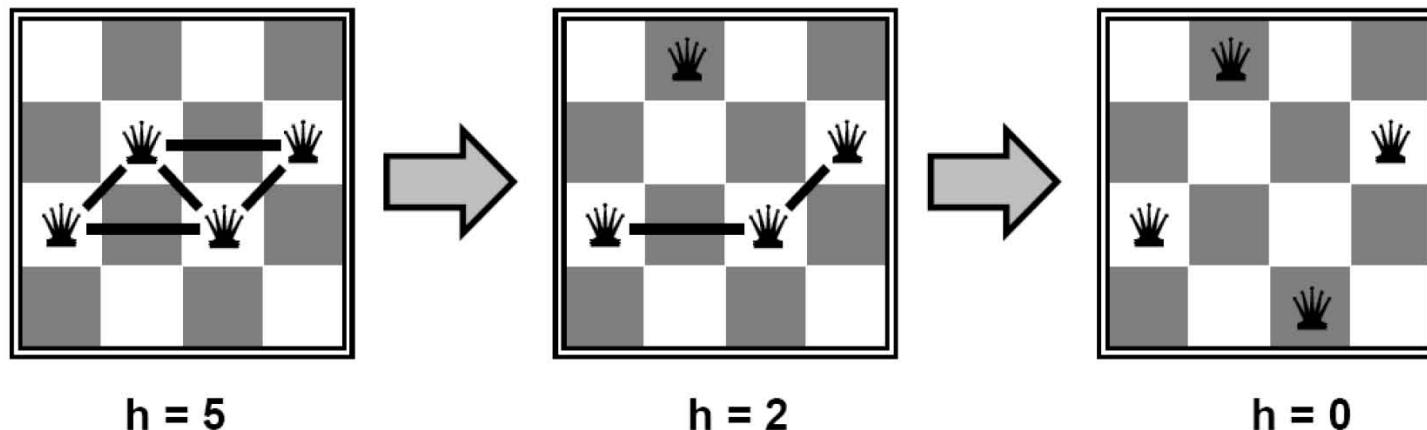
- Is a **configuration** that **satisfies all constraints**

Approach

- **keep a single "current" state** (or a fixed number of them)
- try to improve it by **maximizing a heuristic function**
- use only **“local” improvements**
 - i.e., only **modifies the current state(s)**

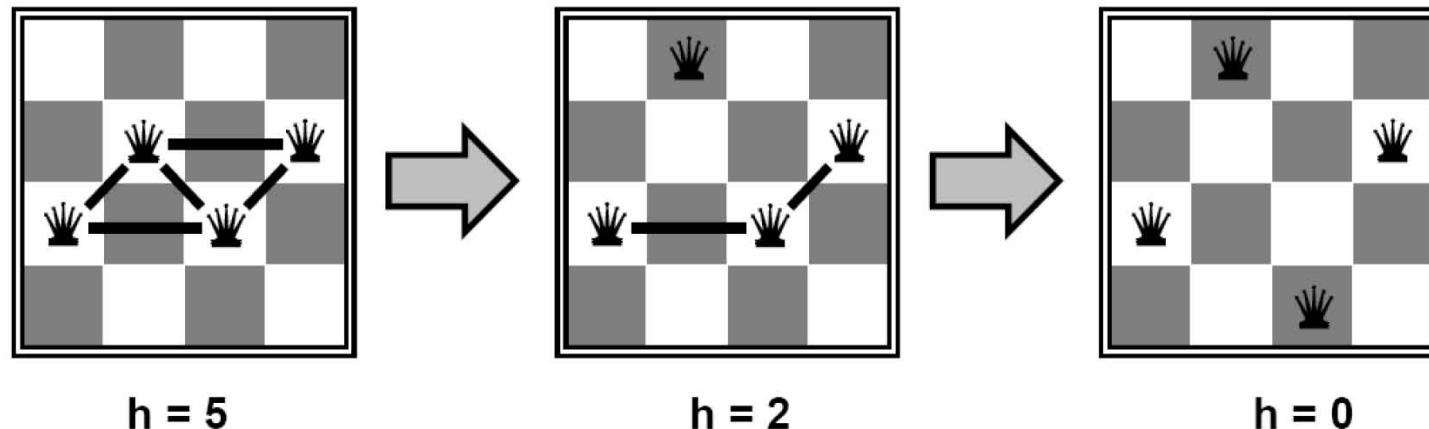
Example: n-queens problem

- Put 4 queens randomly on the chessboard
- move one queen so that it reduces the number of conflicts
- often solves n-queens problems almost instantaneously for very large problems
 - e.g. $n = 1,000,000$ can be easily solved



Example: n-queens problem

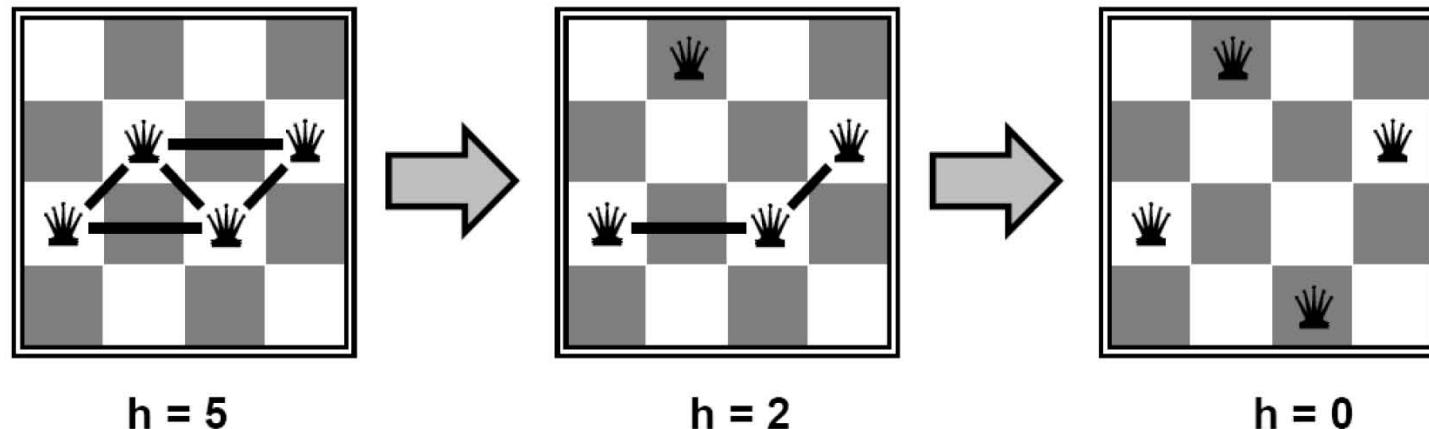
- Put 4 queens randomly on the chessboard
- move one queen so that it reduces the number of conflicts
- often solves n-queens problems almost instantaneously for very large problems
 - e.g. $n = 1,000,000$ can be easily solved



It's like we are minimizing a function that gives the number of conflicts

Example: n-queens problem

- Put 4 queens randomly on the chessboard
- move one queen so that it reduces the number of conflicts
- often solves n-queens problems almost instantaneously for very large problems
 - e.g. $n = 1,000,000$ can be easily solved



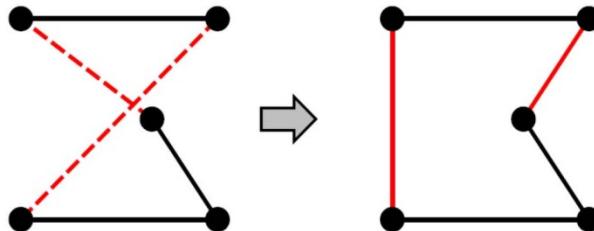
**It's like we are minimizing a function that gives the number of conflicts
Or we are maximizing a function that estimates the value of the state**

Example: Traveling salesperson problem

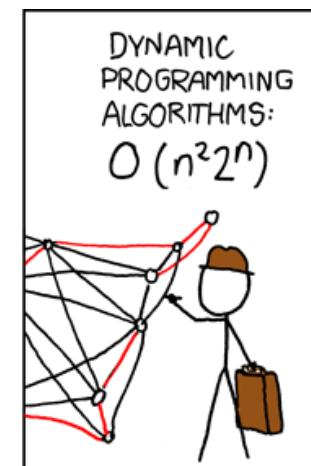
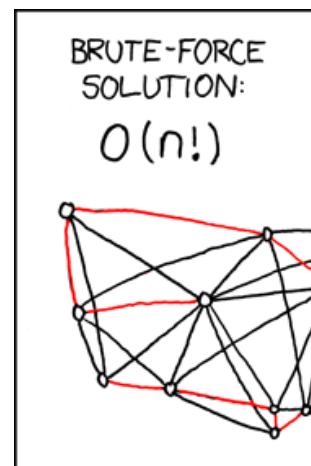


Example: Traveling salesperson problem

- start with a complete tour that passes all cities
- perform pairwise exchanges so that the exchange reduces the cost of the tour
- variants of this approach quickly get close to an optimal solution even for thousands of cities

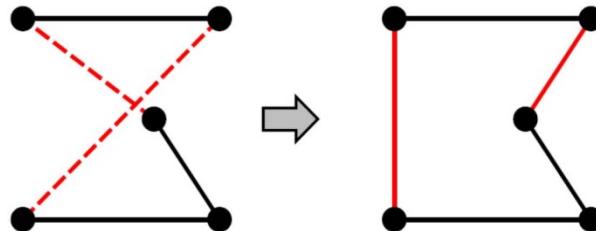


This will be our first challenge!

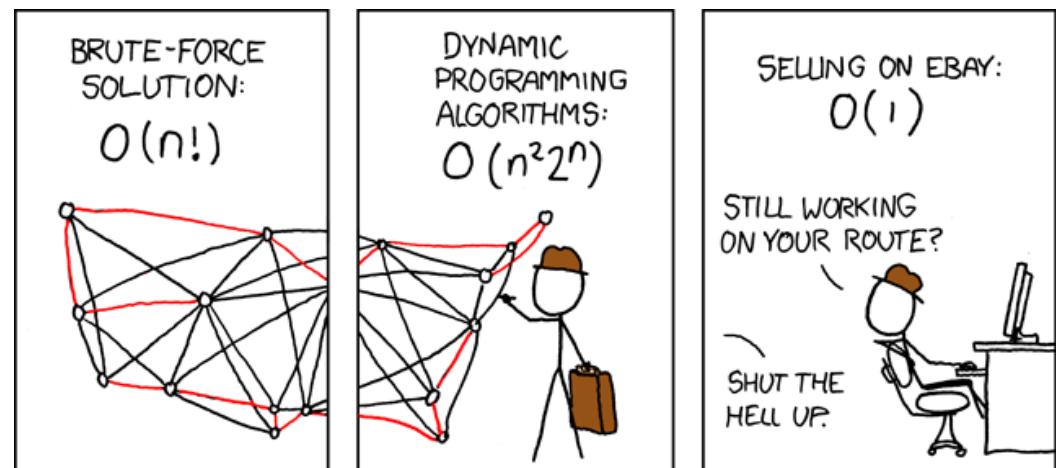


Example: Traveling salesperson problem

- start with a complete tour that passes all cities
- perform pairwise exchanges so that the exchange reduces the cost of the tour
- variants of this approach quickly get close to an optimal solution even for thousands of cities



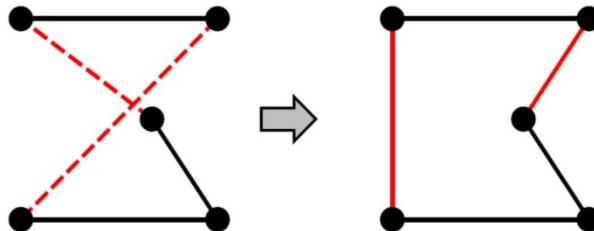
This will be our first challenge!



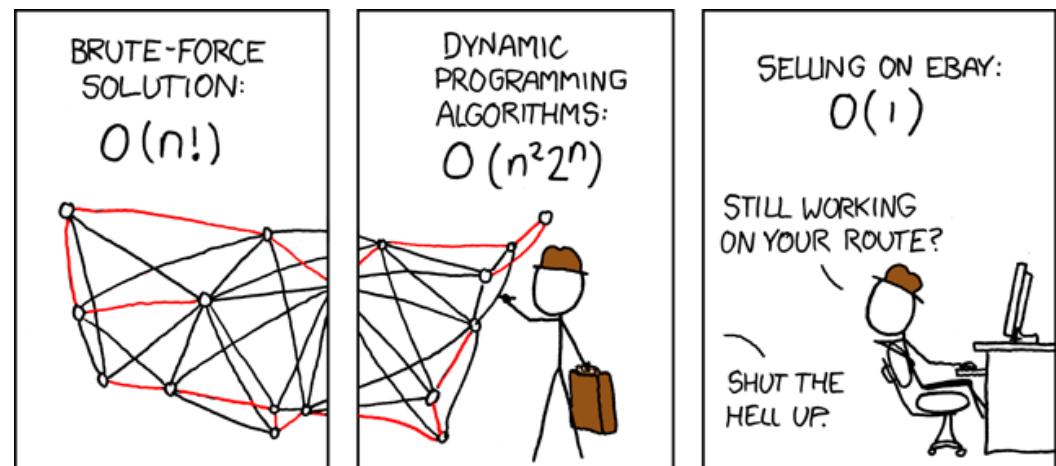
It's like we are minimizing the a function that gives cost of a tour

Example: Traveling salesperson problem

- start with a complete tour that passes all cities
- perform pairwise exchanges so that the exchange reduces the cost of the tour
- variants of this approach quickly get close to an optimal solution even for thousands of cities



This will be our first challenge!



It's like we are minimizing a function that gives cost of a tour

Or like we are maximizing a function that estimates the value of the state

Local search

Advantages

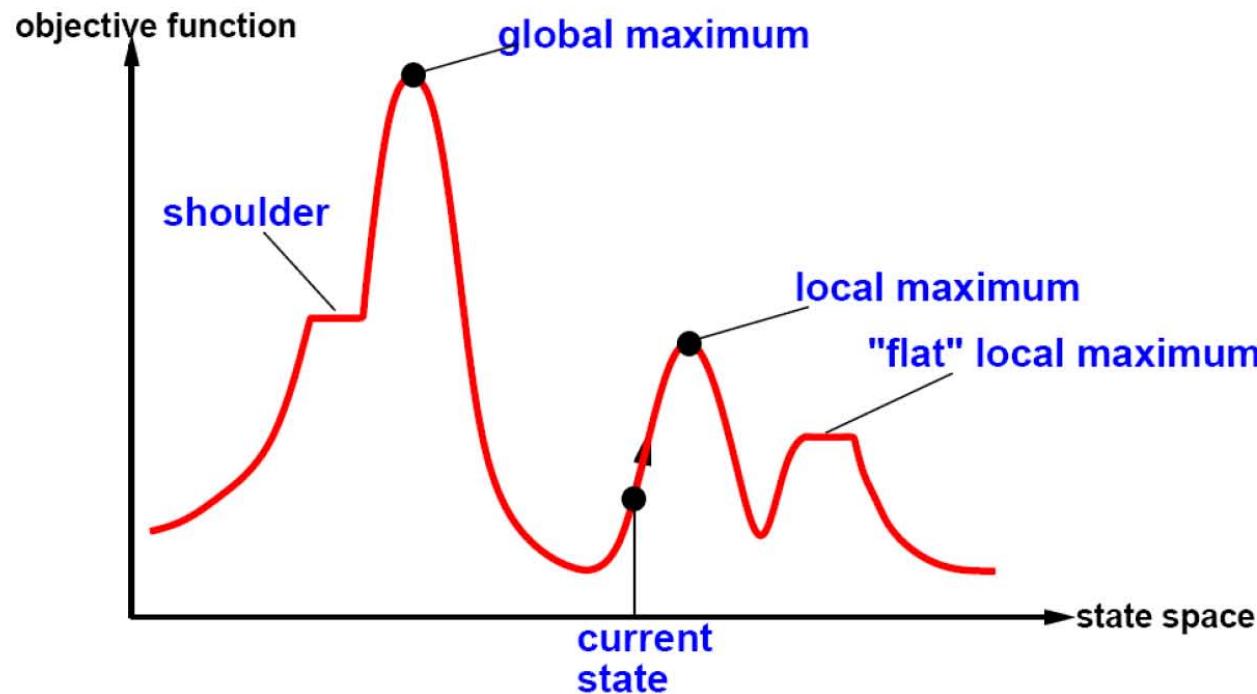
- It uses very **little memory** —> we keep only one or a few states
- often **quickly finds solutions** in large or infinite state spaces

Disadvantages

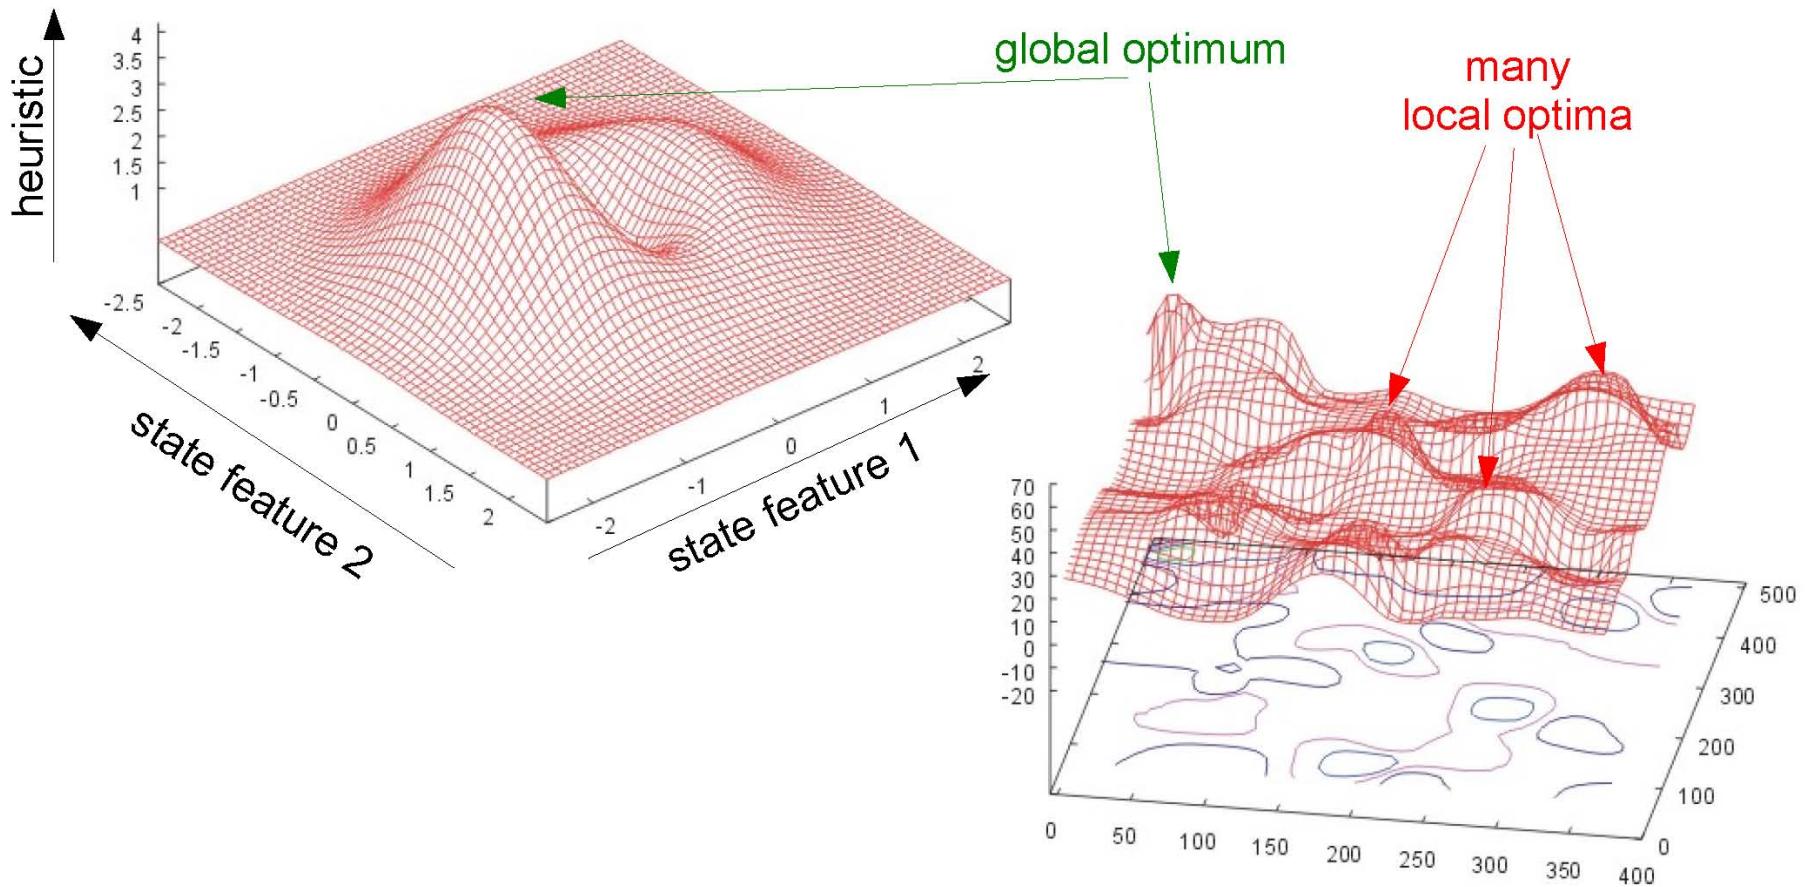
- no guarantees for completeness or optimality

State space landscape

- location: states
- elevation: value of objective function
- problematic regions: local maximum, plateau



Multi-dimensional state landscape



Local search algorithms

- Hill climbing
- Beam search
- Simulated annealing
- Genetic algorithms

Hill Climbing

Hill climbing search (greedy local search)

- Expand the current state → generate all neighbor states
- Neighbor states: states that can be reached within a single transformation of the current state
- Move to the one with the highest evaluation.
- Continue as long as the evaluation score increases.

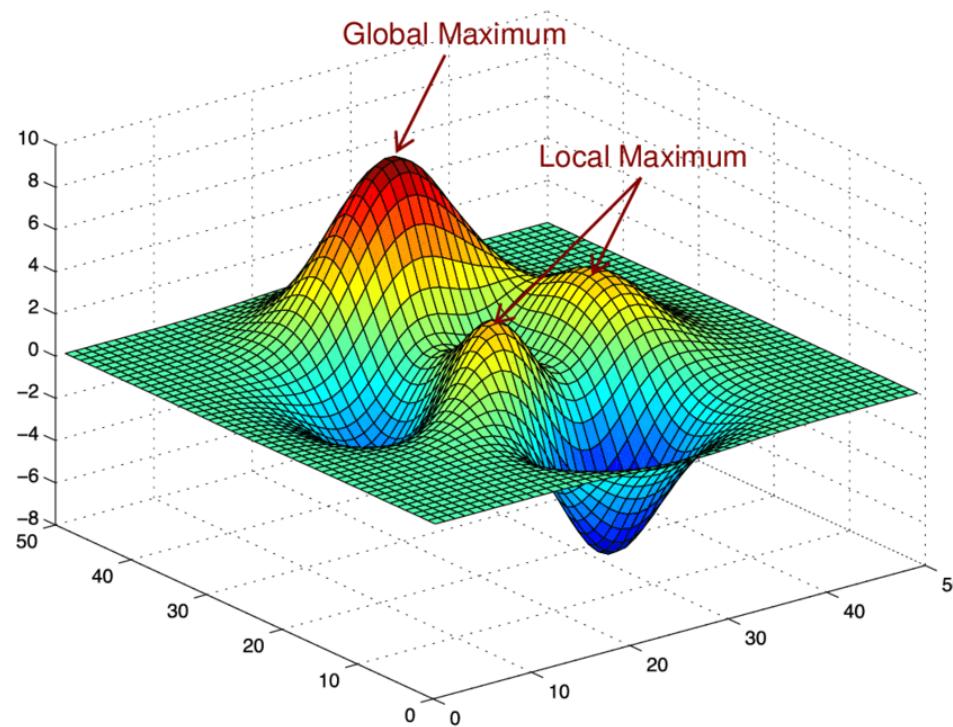


https://c1.staticflickr.com/8/7140/7486564384_695593d27a_b.jpg

Hill climbing search

Main problem: Local optimum

- The algorithm will stop as soon as it is at the top of a hill.
- But it is actually looking for a mountain peak.



https://www.researchgate.net/figure/11-Illustration-of-local-optimum-and-global-optimum_fig16_306558608

8-queens problem

Heuristic h as the evaluation function

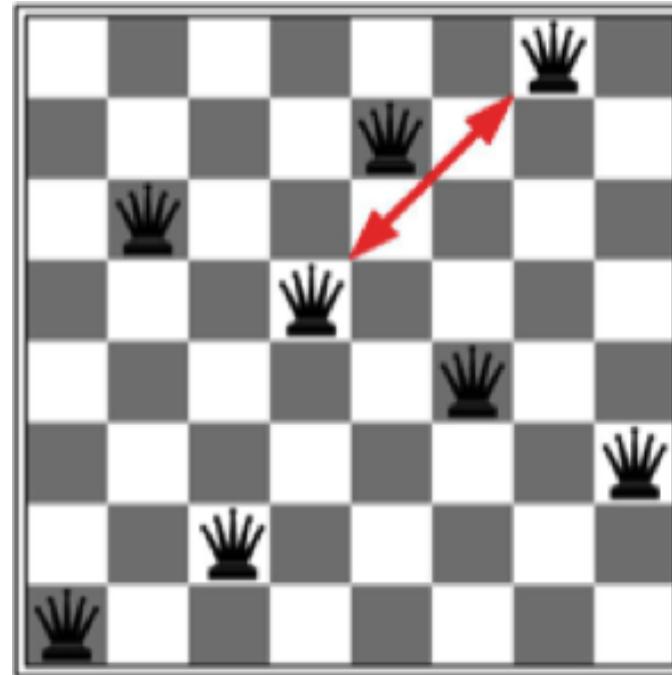
- Number of pairs of queens that attack each other → minimize h
- Example state: $h=17$
- Best neighbor: $h=12$
→ Several possible moves

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	14	13	16	13	16
14	14	17	15	13	16	13	16
17	14	16	18	15	14	16	16
18	14	15	15	15	14	15	16
14	14	13	17	12	14	12	18

Example: 8-queens problem

Stuck in **local optimum** with $h=1$

- No queen can move without decreasing the number of attacked pairs



Randomized hill climbing variants

Random restart hill climbing

- **Different initial states** result in **different local optimum**
→ We should run **several iterations** with **different starting states**

Stochastic hill climbing

- Select the **successor state randomly based on the** probability of each state
- The probability of a state is estimated using the heuristic value for the state

Beam Search

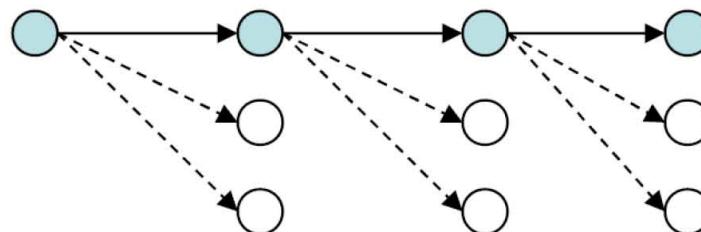
Beam search

Idea:

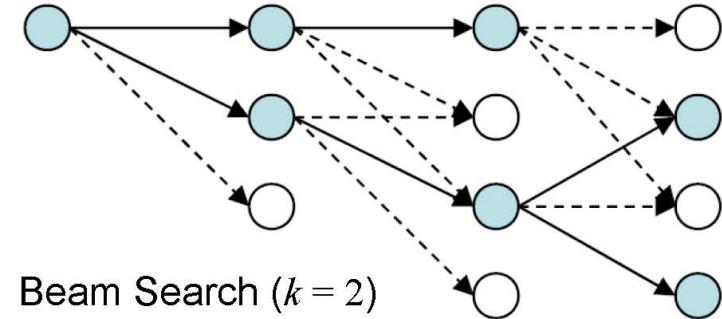
- Keep track of k states rather than just one
- k is called **the beam size**

Algorithm

- Start with **k randomly generated states**
- At each iteration, generate **all successors of all k states**
- If any state **satisfies the constraints**, stop;



Hill-Climbing Search



Beam Search ($k = 2$)

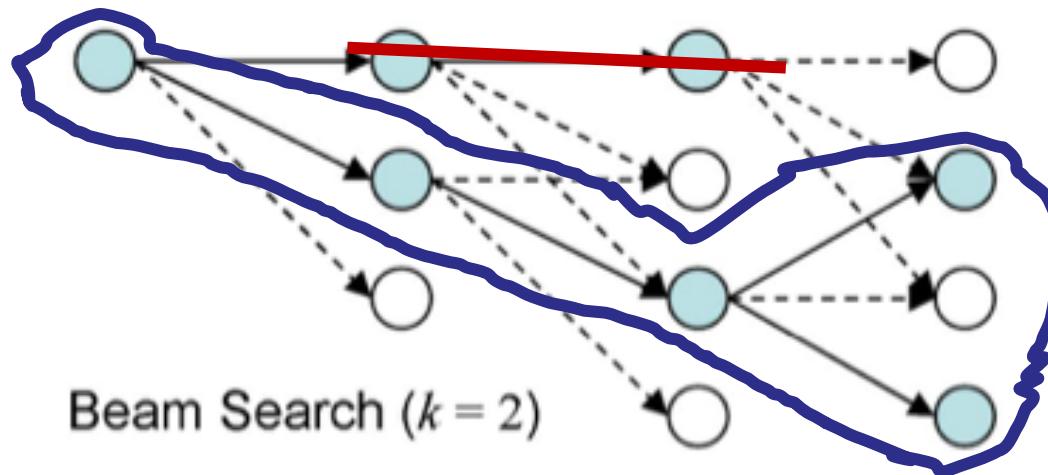
Implementation

- similar to the **tree search algorithms**
 - sort the **queue** by **the heuristic function h** (as in greedy best-first search), but **limit the size of the queue to k**
- expand all nodes in queue simultaneously (not in sorted order)

Beam search

Note

- **beam search** is **different** from **k parallel hill-climbing** searches!
- information from different beams is combined



Beam search

Effectiveness

- suffers from **lack of diversity of the k states**
 - e.g., if one state has better successors than all other states, alternative routes are pruned early
 - **often** not more effective than basic hill climbing

Stochastic beam search

- chooses **k successors at random**
- selection probability of each states depends on heuristic value for the state

Simulated Annealing

Simulated annealing

Algorithm:

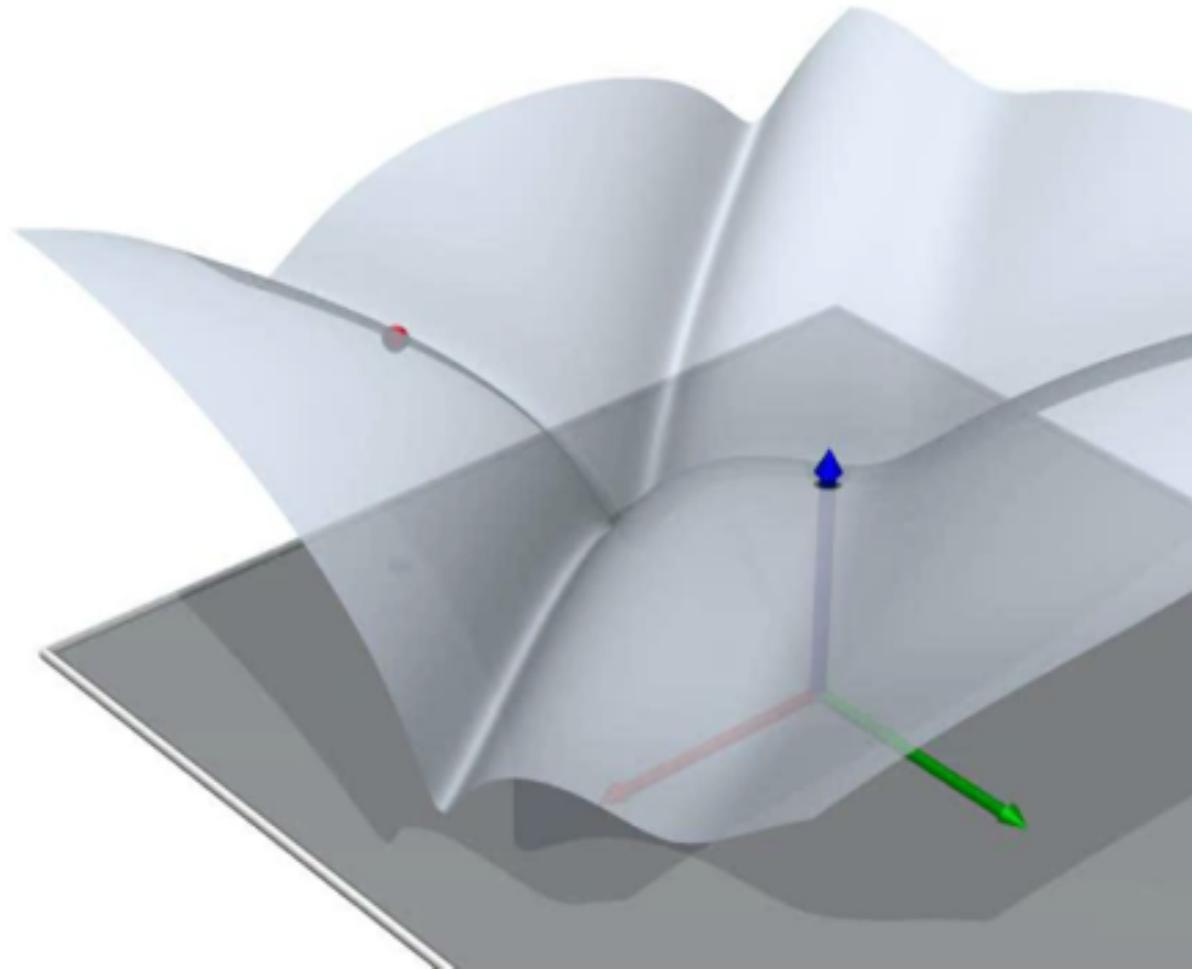
- temporarily accept worse configurations but with decreasing probability
- Combination of **hill climbing** and **random walk** through the state space
 - **escape local maximum** by **allowing some “bad” moves**
 - but gradually **decrease their frequency** (the *temperature*)

Effectiveness

- it can be proven that if **the temperature is lowered slowly enough**, the **probability of converging to a global optimum approaches 1.**

Widely used in VLSI layout, airline scheduling, etc.

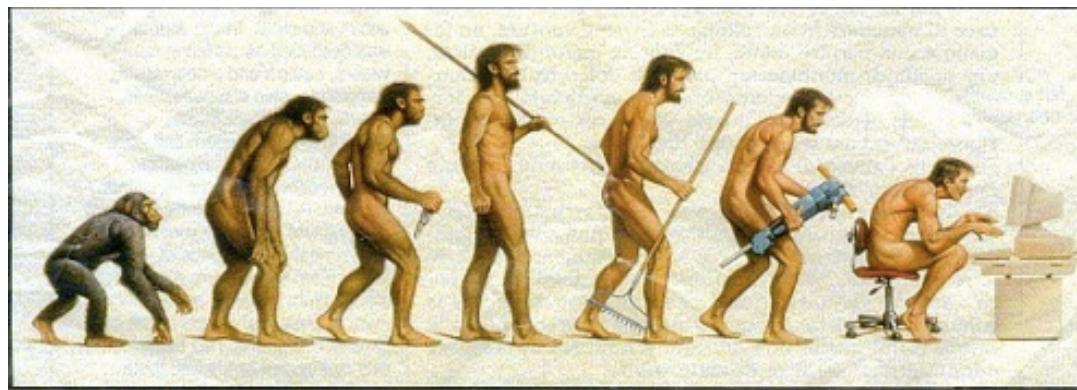
Simulated annealing



Genetic Algorithms

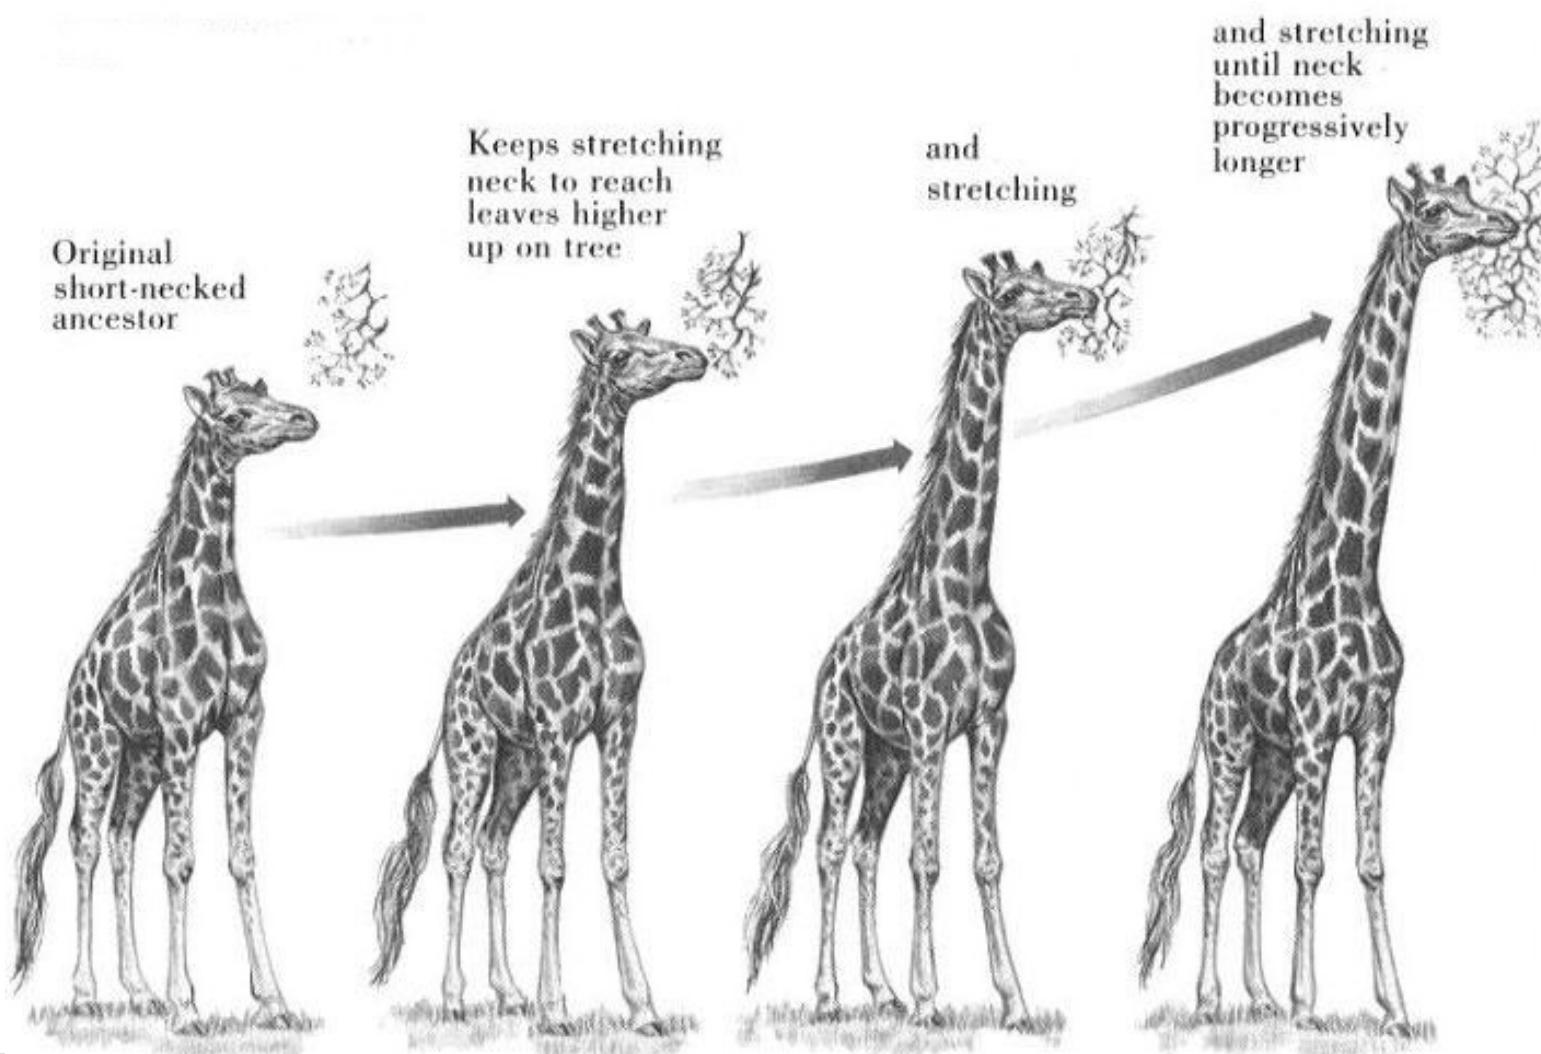
Metaphor: Genetic algorithms

- Analogy with evolutionary theory → Simulation of natural selection
- View **states** as **individuals in a population**
- **Objective function** measures the **fitness of an individual**
- Allow “**fittest**” individuals **to reproduce**
- **Eventually**, the **population** should evolve towards **fitter individuals**.

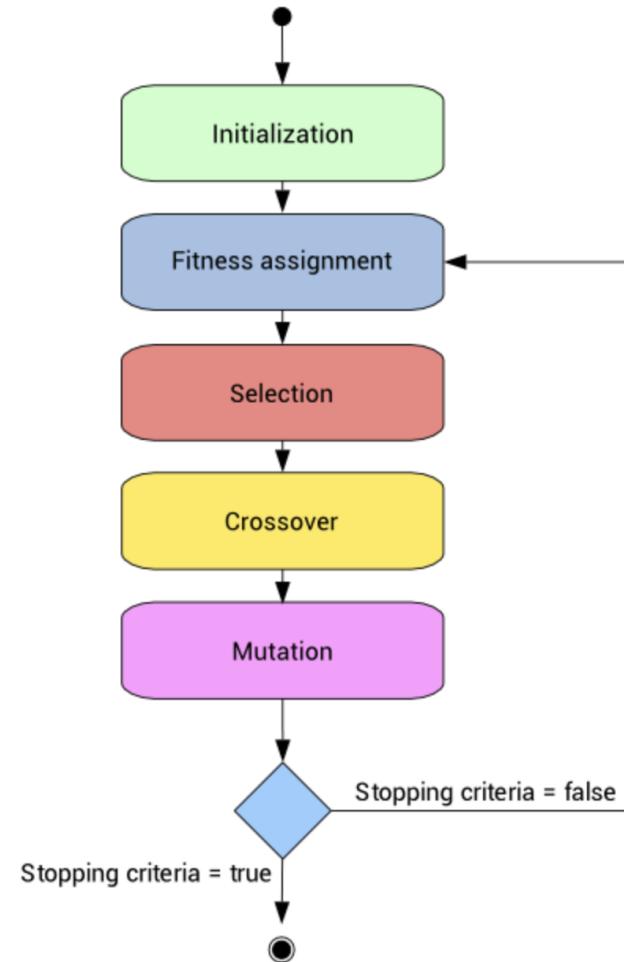


Somewhere, something went terribly wrong

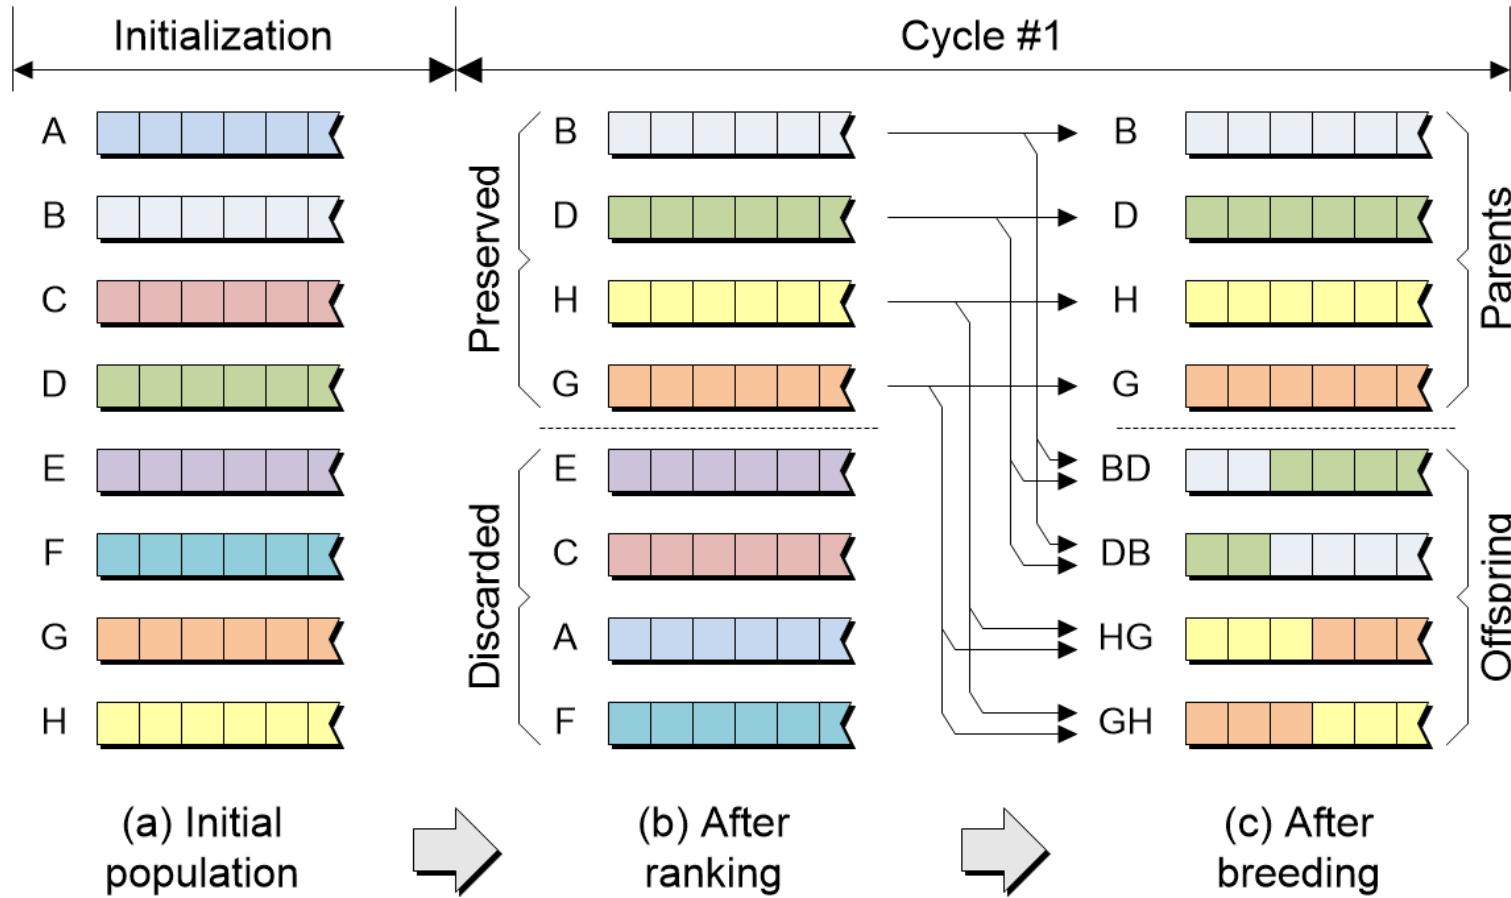
Genetic evolution



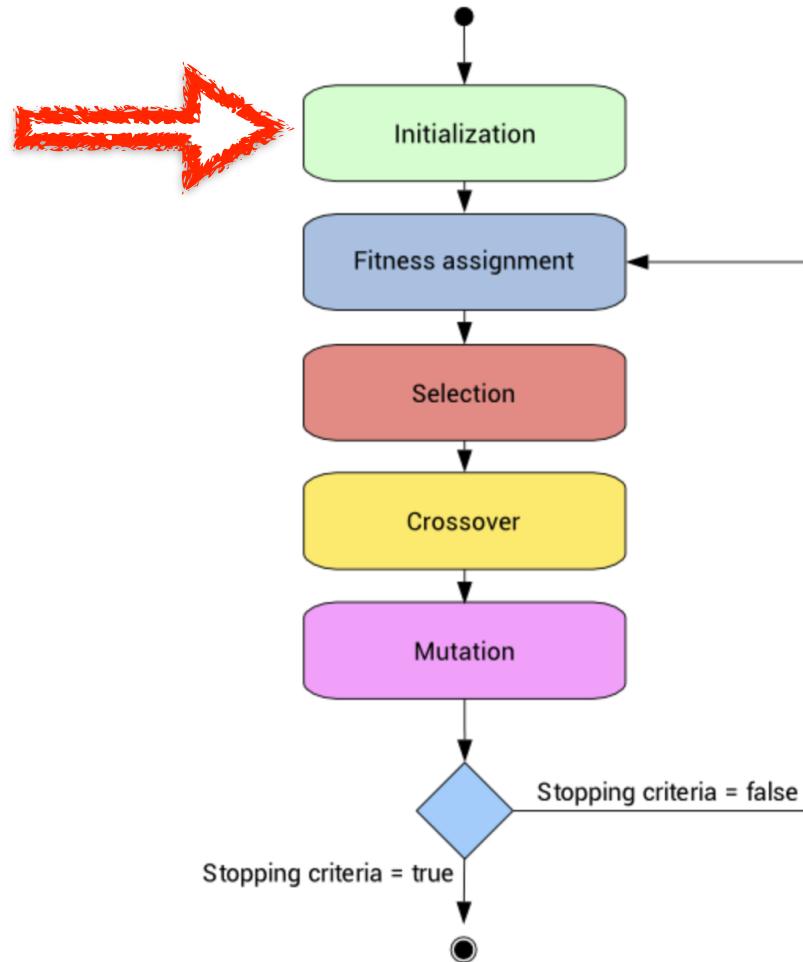
Genetic algorithm



Genetic algorithm



Genetic algorithm



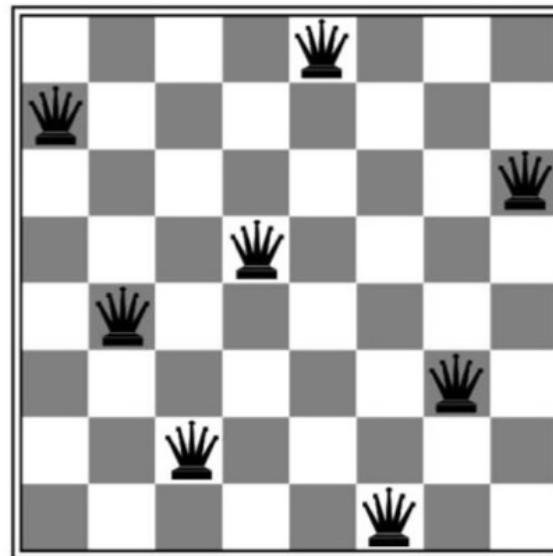
State = problem configuration

The problem needs to be represented by a configuration.

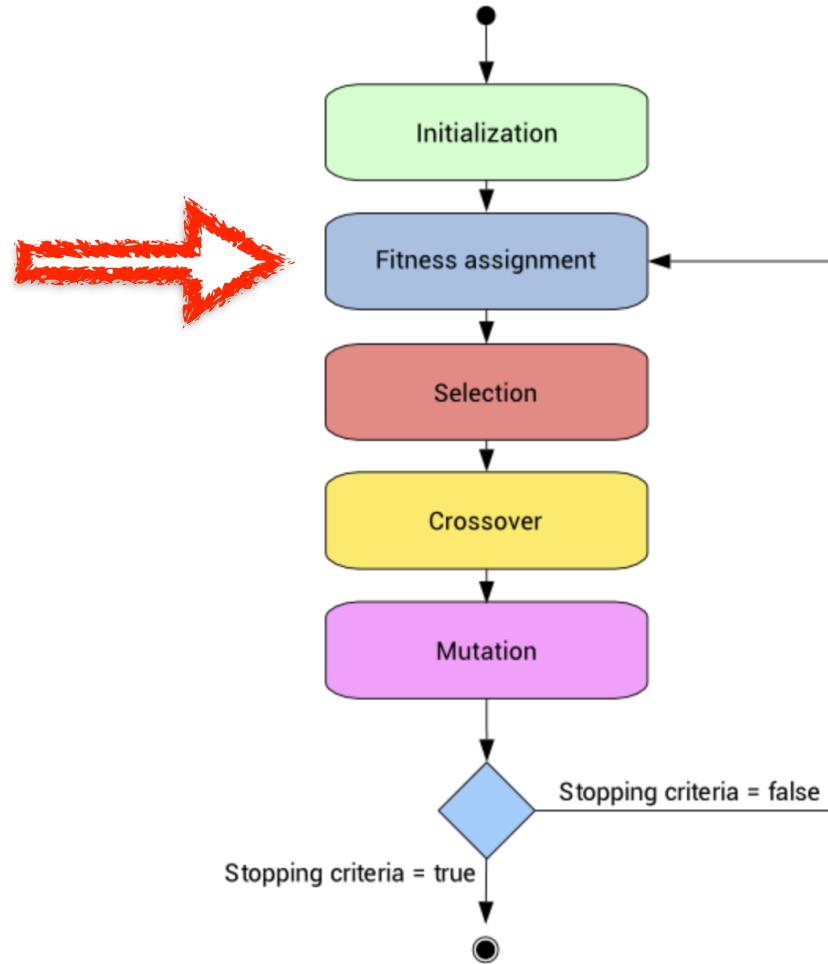
$$X = \boxed{1 \ 0 \ 0 \ 1 \ 1 \ 0 \ 0 \ 1}$$

8 queen representation

$x = 74258136$



Genetic algorithm



Fitness assignment and selection

- ◊ Start with a **random population** as the initial state
- ◊ Apply an **evaluation function (fitness)**
 - ◊ should assign **higher values** to better states
 - Number of **non-attacking** pairs of queens
- ◊ **determine selection probability** according to fitness

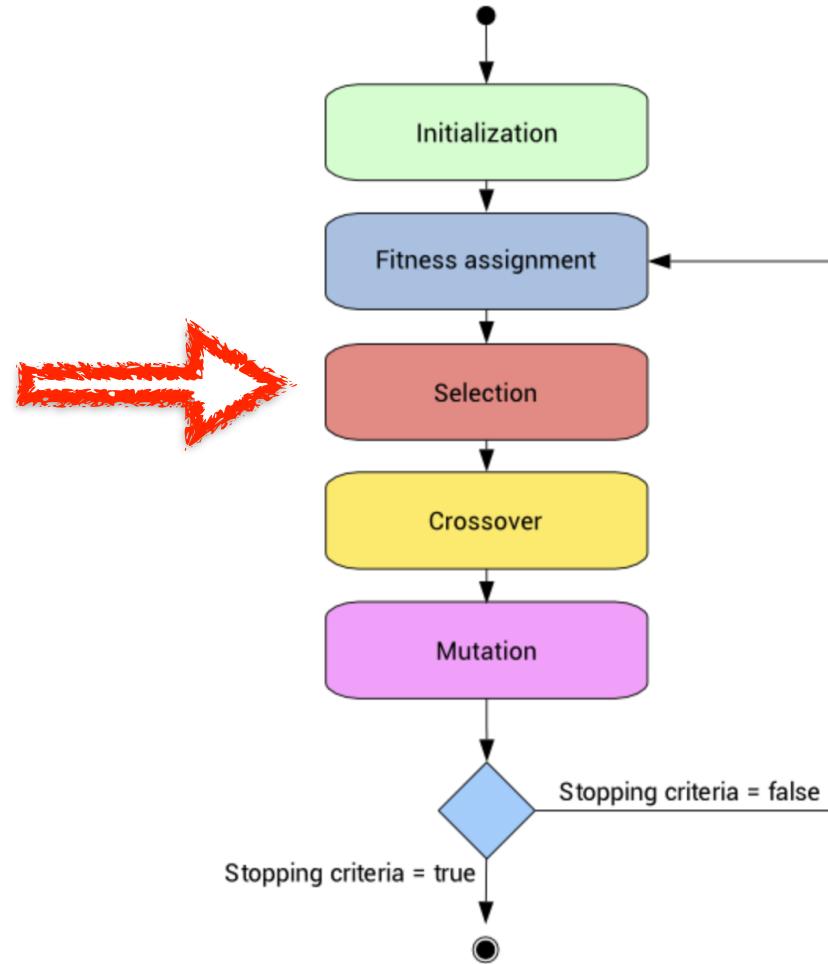
24748552 → 24 → 31%

32752411 → 23 → 29%

24415124 → 20 → 26%

32543213 → 11 → 14%

Genetic algorithm



Selection

- ◊ randomly select states with respect to their selection probability
- ◊ selected states will be parents for the new generation of the population
- ◊ Parameter: **selection rate R** = what portion of population can survive?

24748552 → 24 → 31%

32752411

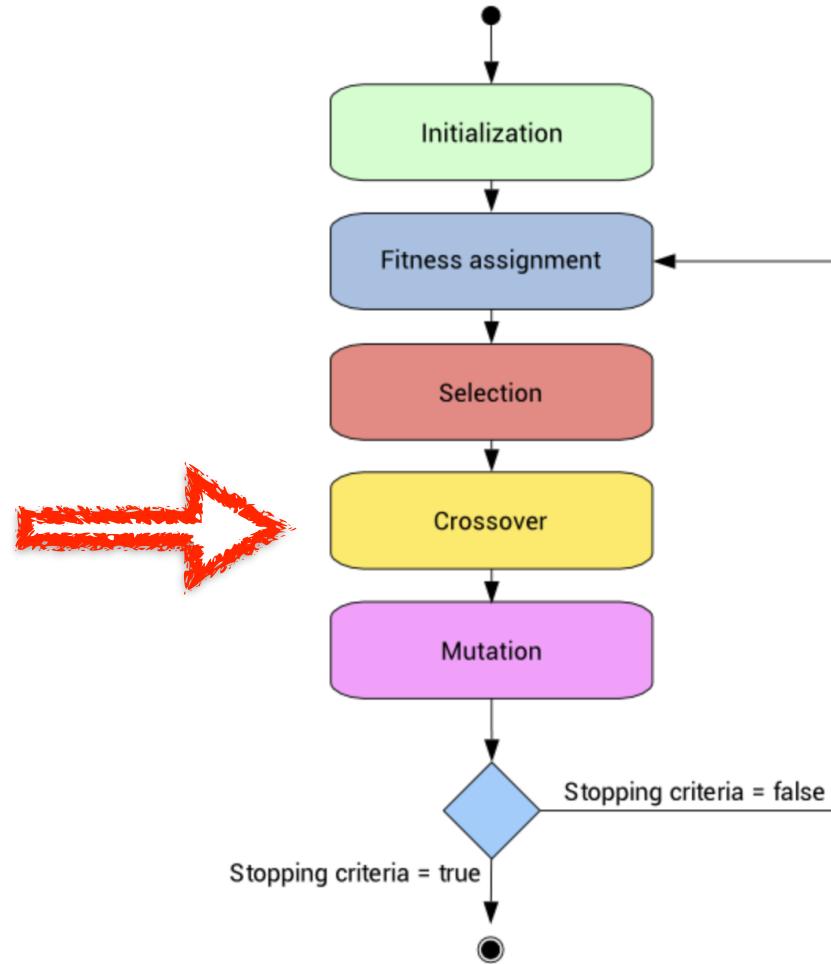
32752411 → 23 → 29%

24748552

24415124 → 20 → 26%

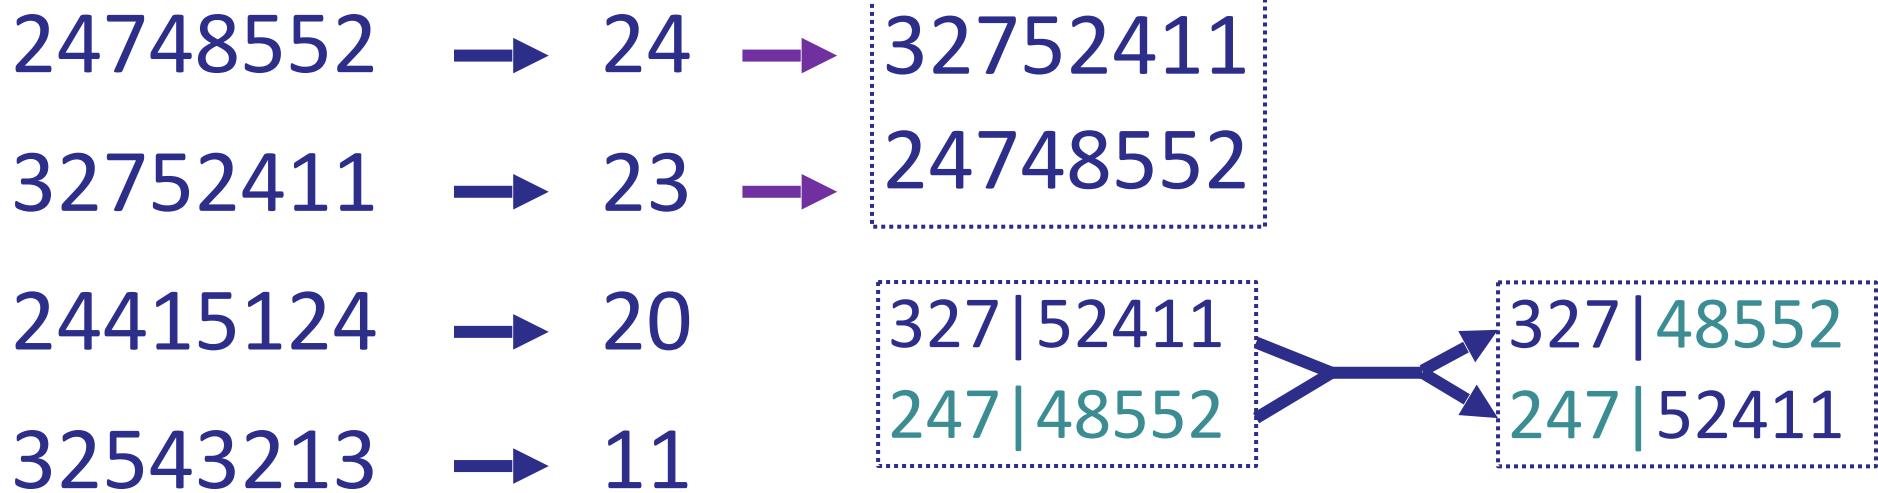
32543213 → 11 → 14%

Genetic algorithm

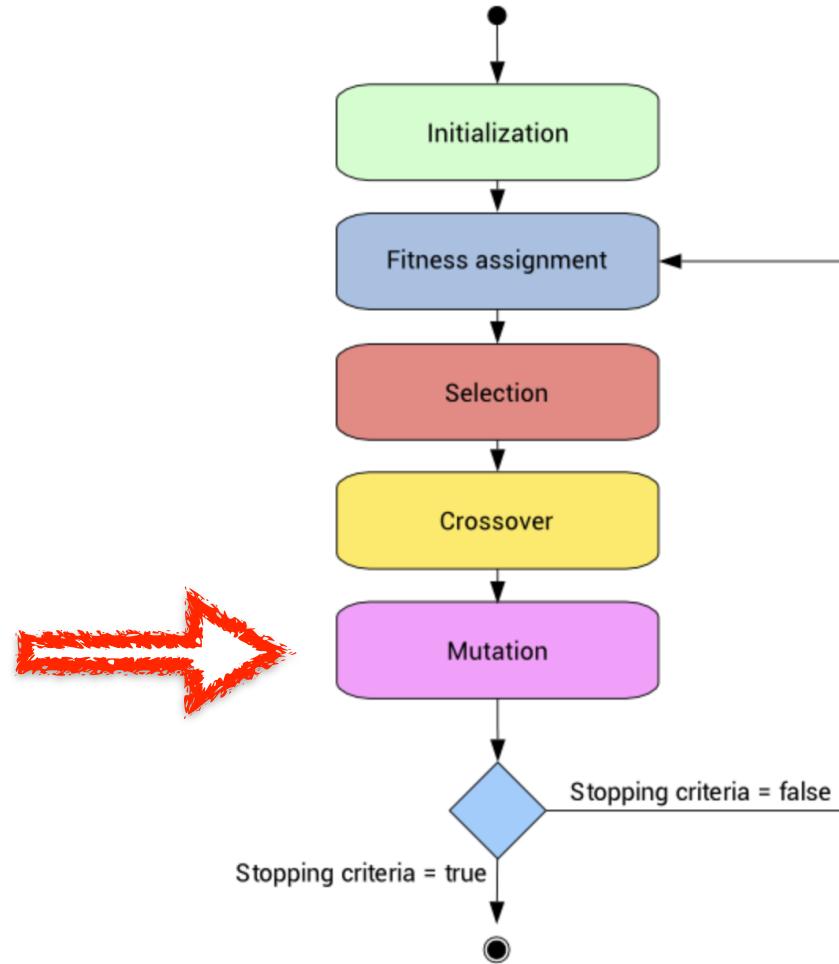


Cross-over

- take two parent states
- randomly determine the cross-over point
- cut parents at cross-over point
- recombine the pieces

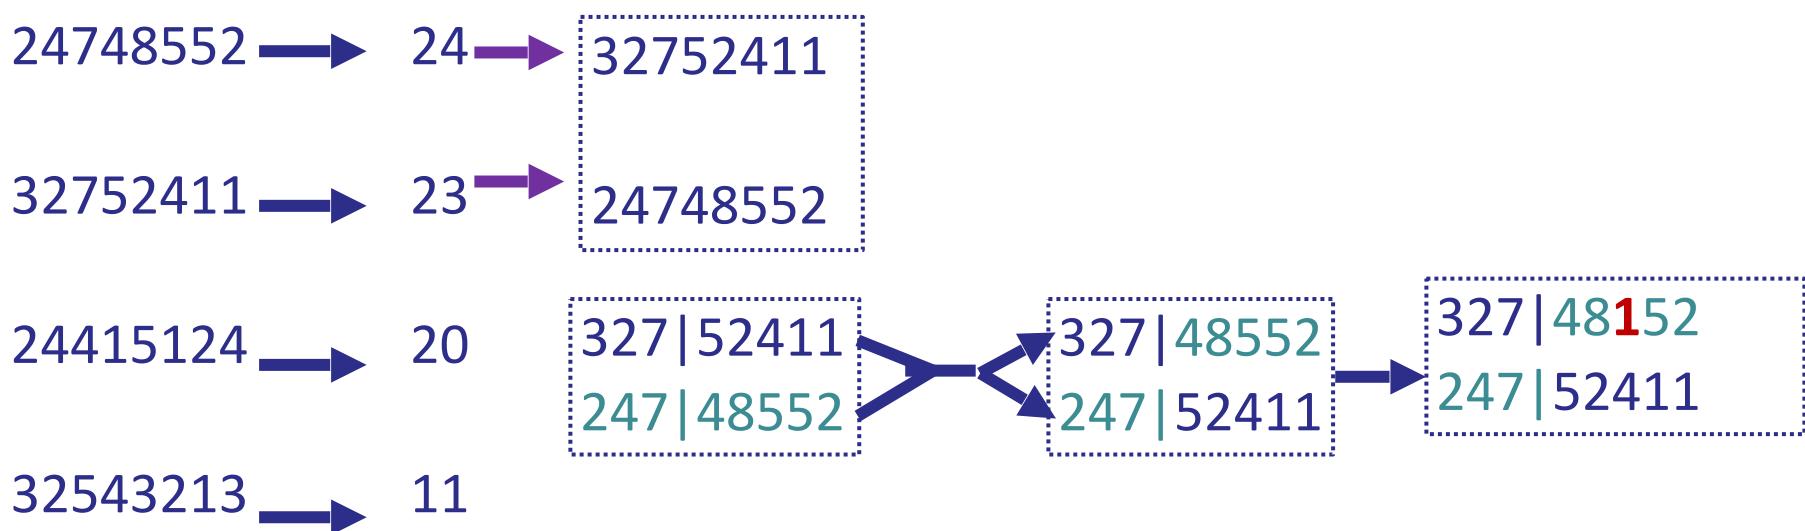


Genetic algorithm

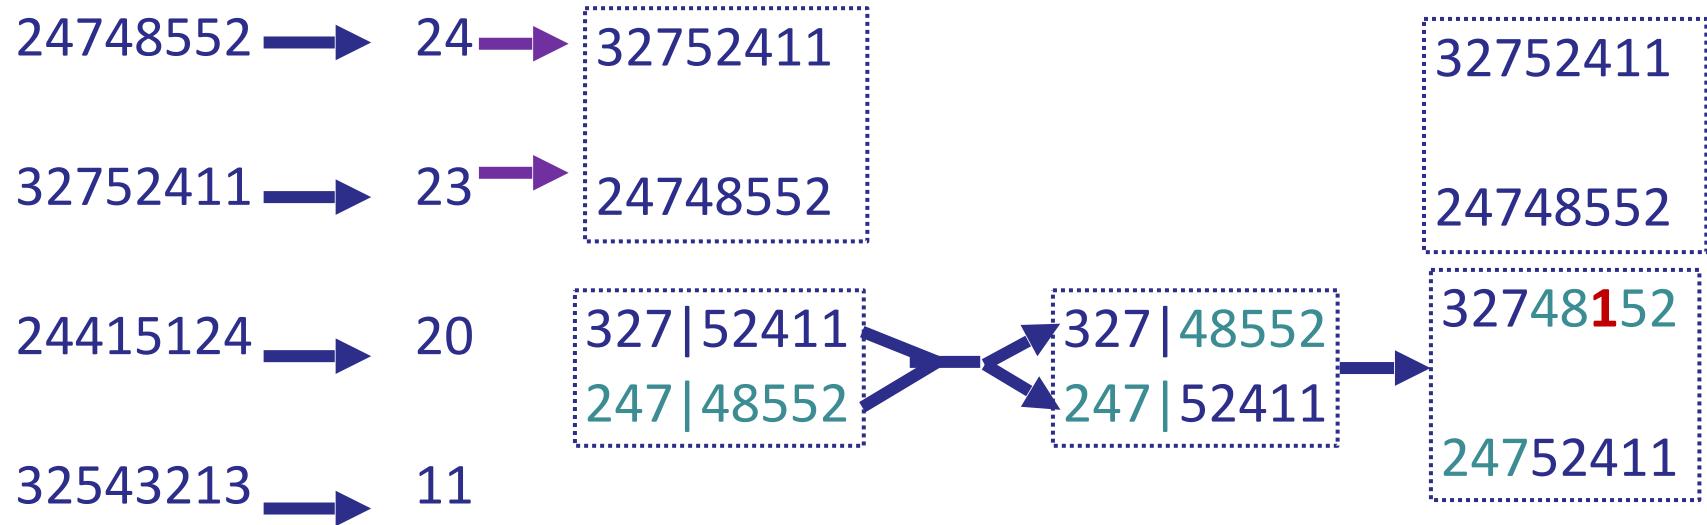


Mutation

- Random change of **n** elements in the output states of cross over
 - corresponds loosely to random walks
 - **Parameter: Mutation rate μ : what portion of elements in states should mutate?**



New population

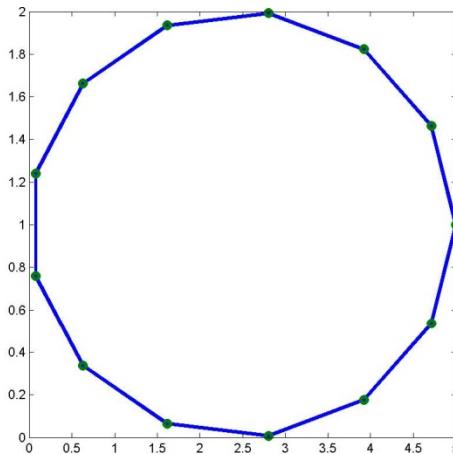


Genetic algorithms

Properties

- **easy to implement** and **to explain** general **optimization algorithm**
- **performs well**
 - **unclear under which conditions** they work well
 - other randomized algorithms perform equally well (or better)
- **Applied to optimization problems:**
 - circuit layout,
 - traveling salesman problem (TSP),
 - data clustering and mining,
 - image processing,
 - medical science

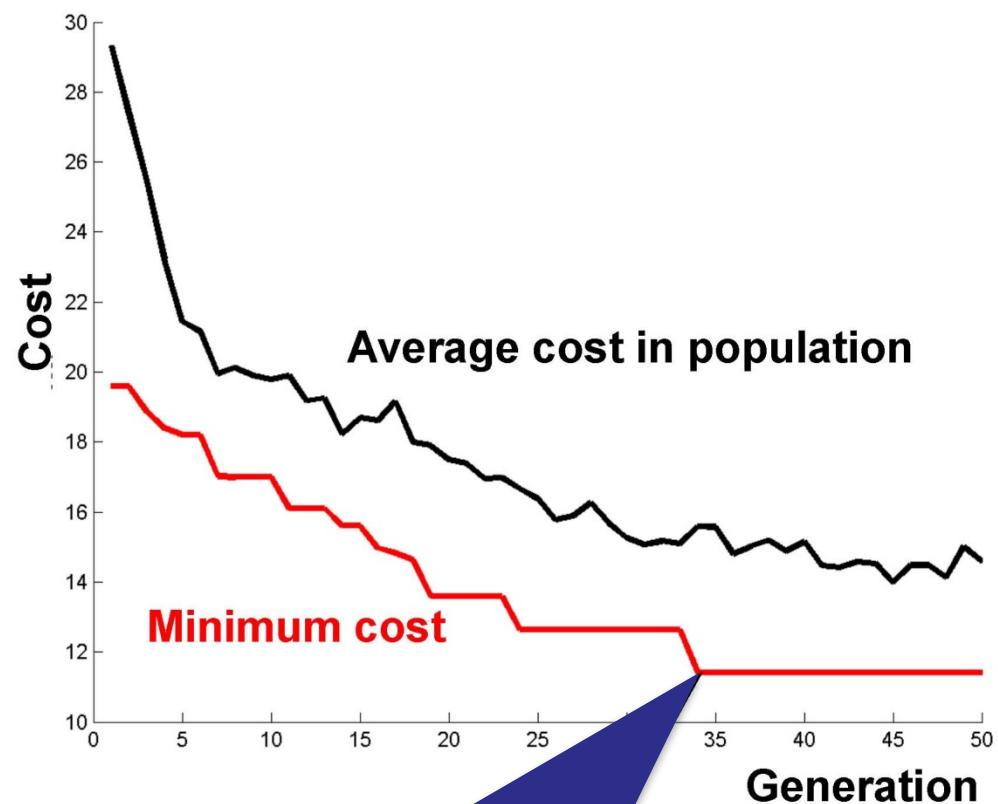
Example: Traveling salesperson problem



$N = 13$ cities

$P = 100$ states in population

$\mu = 4\%$ mutation rate
 $R = 50\%$ selection rate



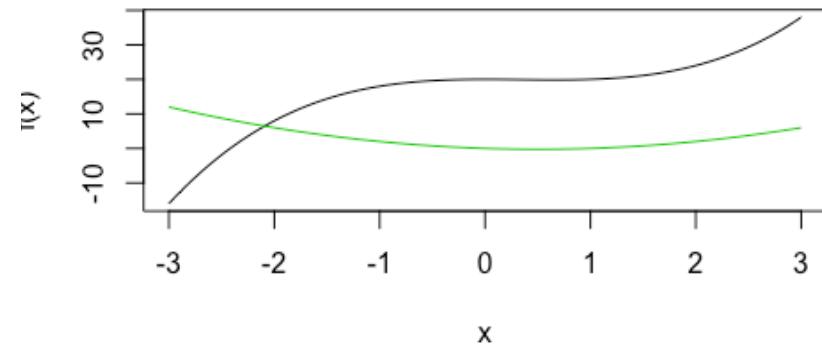
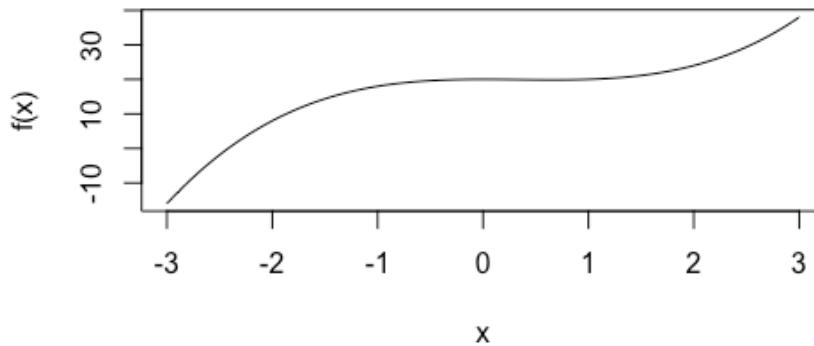
Optimal solution reached at generation 35

Local search in continuous spaces

In many real-world problems, the **state space is continuous**.

How to define neighboring states?

- **Discretize** the state space
 - e.g. assume only **n** different positions of a steering wheel or a gas pedal
- **Gradient Ascent**
 - **hill-climbing** using **the gradient of the objective function f.**
 - What is the **gradient? Vector of partial derivations.**



Summary

- Local search: only focuses on the final configuration (not on the path)
 - Methods
 - Hill climbing
 - Beam search
 - Simulated annealing
 - Genetic algorithm
- find good solutions for many practical problems
- require a careful representation of the problem
- State representation
 - Evaluation function

Readings

Mandatory

- Russell & Norvig, Section 4: *Beyond Classical Search*
 - 4.1 *Local search algorithms and optimization problems*, p. 120-129

Optional

- 4.2 and 4.3
- Rest of chapter 4

